

Principles of Computer Systems Design - Assignment 1

Daniel Egeberg & Søren Dahlgaard

December 5, 2012

1 Exercises

Question 1

Our scheme is very simple: First calculate which machine to contact, and what the memory address is on that machine. Second contact the machine over ethernet (ie. some local LAN).

We make the assumption, that we are aware of how much memory each individual machine has, when creating the scheme. Two cases then present themselves:

1. Each machine has the same amount of memory it is very easy to calculate the machine and address using $a \bmod m$ and $\lfloor a/m \rfloor$, where a is the address, and m is the amount of memory.
2. If this is not the case, it is slightly more complicated. One way to handle this is to keep an array in the translation code, which contains the accumulated sizes, and then doing a binary search on this array to find the machine. Eg. if machine 0 and 1 have 2 GB of memory and machine 2 and 3 have 4 GB of memory, the array would contain the values $[0, 2 \cdot 10^9, 4 \cdot 10^9, 8 \cdot 10^9]$. We then simply do a binary search for the biggest number in the array, that is smaller than the sought address a . Then calculating the local address on the machine is trivial. This adds a log factor of the number of machines, but this number will never be big enough to be a dominating factor.

To have the API look as similar to the one for local memory, we limit READ and WRITE to only accessing one byte at a time. See the pseudocode below. The pseudocode is for the one-size memory version, but the other one is almost just as simple. Assume m to be some global constant.

READ(a, p)

- 1 $d = a \bmod m$
- 2 $a2 = \lfloor a/m \rfloor$
- 3 $p = \text{SEND}(d, \text{READ}, a2)$

Here SEND is some generic network procedure, that sends the command READ to the machine at d with parameters $a2$. We assume p to be some reference to store the received data in. The pseudocode is the same for WRITE except the

SEND call will send the command WRITE instead and takes an extra parameter – the value to write.

Read/write operations to regular main memory is clearly atomic, as it is performed with a single hardware instruction – although the electric signal could end up being corrupted, this is *extremely* unlikely. Since our method only writes one byte at a time, and this is done using one instruction in the receiving machine, it is clear that our method is atomic as well. We do this because we want the API to be as similar to performing the operations on regular main memory as possible.

Question 2

We have filled the table using values from [3], [2] and [1]. Prices and capacities have been found using www.edbpriser.dk. Not that access time is for everything - this includes rotational latency and seektime for HDDs. We have not split into read and write, but rather split into Access time and transfer rate, as we believe this to be more relevant. For reliability we considered how long a unit is expected to last. For instance, the moving parts of the HDD is subject to wear and tear and thus will wear down. The transfer rates have not been normalized, but are with different units – We assume that $1\text{Word} = 4\text{Bytes} = 32\text{bits}$. For sizes we have listed commonly available size only. For example, even though SSDs of more than 2 TB are available, these are not common at all.

	HDD	SSD	RAM
Access time (total)	12ms	0.1ms	33.75ns
Transfer rate	1023 Mbit/s	100 to 600 MB/s	1600MWord/s
Capacity (common)	Up to 4TB	Up to 512GB	Up to 32GB
Cost per GB	0.4 DKK	5 DKK	32 DKK
Reliability	Moving parts subject to wear and tear	Only allows limited number of writes before fail (typically many years under normal use)	Good
Power consumption	2 to 5 W	1/3 of HDD	5 to 15 W

Which kind of storage to use for our abstraction depends a lot on how we wish to use the abstraction. We see it as a huge memory abstraction used in a LAN setting, where the ethernet connection is very fast. In this case HDDs would certainly incur a too big latency overhead to be practical. Also, since writes to disks like HDDs and SSDs are page based we cannot ensure the atomicity we wish to keep our abstraction as close to working with normal memory as we can. Thus we would only use RAM.

Question 3

a)

Concurrency does not have a big impact on latency. This is because each instruction is not performed faster, but rather several instructions are performed at once. This gives bigger throughput, but not better latency. In fact the overhead incurred by the concurrent requests is more likely to have a slight negative effect due to more bookkeeping. Also some requests may need to acquire locks to access shared resources, which increases the latency of the individual request.

b)

Batching and Dallying are two entirely different things. The only thing they have in common is that they can be used to fight bottlenecks. Batching is when you group several requests together before sending them, to avoid the setup overhead. Dallying is when you don't send a request at all because it might not be needed.

An example of batching could be retrieving multiple records from a database at once, by eg. performing an SQL select of all items with id between 10 and 20, rather than retrieving one record at a time. An example of dallying could be in cache, where a memory write is not written to main memory immediately when issued, but is delayed until the cache line has to be replaced.

c)

Caching is clearly an example of fast path optimization. When working with memory it is very common to work with the same area a lot, eg. reading a value adding something to it, and storing it again or searching through an array. A cache optimizes for this, and thus the common case by providing a fast path to this memory once it has been referenced once.

Question 4

a)

Assuming that we are testing the L1 cache, we need to know the size of a cache line. We could just read the same memory address a lot of times, but this is probably not too good, as we will run into too many different optimizations in both software and hardware. Rather, we want to read the same few values. Since we cannot decide how cache lines are replaced it is very hard to utilize the entire cache to be sure that we avoid such optimizations. One way to test the bandwidth would be to read some 64 consecutive bytes (the size of a typical cache line) into the cache and then access these over and over. We must of course turn off all compiler optimizations before doing this.

By performing say, 2^{30} reads we can see how long time it takes to do this, and obtain how many seconds are needed to read one GB. This does, however not take advantage of any parallelism in the cache, so reading words or dwords may be the better option for real effective bandwidth.

We use cache reads, as these cause less delay than reads from an instruction cache, but more than a write miss (due to dallying). This gives a good

midway.

b)

It is quite to ensure, that all reads are cache misses. This depends heavily on the replacement strategy used by the CPU cache, as well as which type of cache it is (direct-mapped, set-associative, full associative, etc.). One way to completely avoid this issue, is to only access one element per possible cache line, flushing the cache when the CPU begins. If we assume a 64 byte cache line size, we could access every 64th byte of the memory. This way, we can probably not perform 2^{30} reads unless we have a lot of memory. Instead, we could read some 100 MB and measure the time for these. To get the bandwidth.

2 Implementation

Question 1

Which RPC style to use depends a lot on what the key-value store is supposed to be used for. If it is to be used in potentially life critical operations we would need to use exactly-once or in some cases at-least once. For a simple web service it would probably be sufficient to use an at-most once policy. Then the user can handle the case of a timeout himself and some times he might have to resend a request if it is of high importance.

Question 2

The main parts of our implementation are described below:

`StoreImpl`

Our Store implementation is just a wrapper on top of the provided `MemoryMappedFile` using a `RandomAccessFile`.

`IndexImpl`

Our Index implementation consists of two parts: A `TreeMap` that contains the position and size of each key in the store, and a list of free blocks of memory in the file. This way we get a `malloc`-like functionality on the store, such that we can reuse space when keys are removed, by adding the block back into the list, possibly merging it with existing ones. All operations are simply done by looking the key up in the `TreeMap` and then either reading or deleting the position in the store. All the operations are without any form of synchronization, as we handle this in the `KeyValueBaseImpl`.

`KeyValueBaseImpl`

The key-value implementation is really just a glorified wrapper around `IndexImpl`, which also handles synchronization, by using a `ReentrantReadWriteLock`. We do not use per-key locks, as we like the `KeyValueBaseImpl` to handle locks. It would, however, be easy to change this and let the Index handle a lock per key. This, does however add some requirements on `bulkPut` and `atomicScan`, as these would have to acquire all locks in the beginning to ensure atomicity. This could potentially incur deadlocks, unless keys are acquired in a controlled fashion (ie. increasing

by keys). The `init` function of the class reads in the keys one by one and adds them via the index. The only thing we need to consider here, is that a key may span multiple lines, but since these are sorted it is trivial to handle.

We additionally had to create a class `KeyValueBaseService`, which wraps `KeyValueBase`, to make it work with the web service and JAX-WS. This class encapsulates a lot of types because interfaces and collections can't be used with JAX-WS. Instead we use plain arrays for return types in scan operations and wrap the input types in our own classes. All the packing in and out of wrappers is done by `KeyValueBaseService`, to keep the interface of the `KeyValueBase` consistent with the assignment.

Question 3

Question 4

The parameters are discussed in the following list. We only consider the parameters noted in the assignment.

Number of Clients influence the throughput greatly, as the web service is able to handle concurrent requests. Note however, that it might actually decrease both throughput and increase latency due to the mix of operations (see below).

The hardware greatly influences latency, as it decides how much of the memory-mapped file can remain in the main memory, and thus how long it takes to read/write requests.

The size of the dataset influences latency in the same way as the hardware does. If the dataset is huge – less of it can be in main memory, and thus response time will increase.

Mix of operations Because we employ CREW (concurrent read, exclusive write) locking, write operations greatly increase latency of read operations and decreases the throughput.

Question 5

Question 6

Question 7

Question 8

We would definitely like to perform some tests on a computer with much less/more memory to see the latency de/increases we expect from the memory mapped file response times. The same result should be obtainable by using a smaller/bigger dataset.

The experiment from question 5 also only uses **read** operations, which are by far the fastest. We would like to see the effect of adding **write** operations to the mix – eg. having every n th operation as a write for different n , to see how much concurrency we can get out of the reads while they have to wait once in a

while. This should decrease throughput, and increase the latency of *some* of the read operations.

Question 9

As already noted, both methods had to be fitted heavily to work with JAX-WS. This means, that the interface for the end-user (not the `KeyValueBaseImpl`) has changed a bit to allow the user to pass a collection over JAX-WS. Since we don't use per-key locks, `atomicScan` simply acquires the read lock and then calls `scan`. If one read fails the user will get an exception, and thus the entire procedure is atomic.

Our `bulkPut` is implemented by acquiring the write lock and then inserting or updating the entries. While doing so, it keeps a log of all the changes it has made, and tries to restore these if an update fails. This almost ensures atomicity, but has a few flaws. First of all, the log is kept in memory, so if one tries to update the entire store, the log will be too big, and the system is likely to crash. Also, if one of the rollback writes fail, we do nothing, as this is a slippery slope. In the general case it should work well, though.

References

- [1] Tom's Hardware, *Component Power Requirements*, <http://www.tomshardware.com/reviews/power-saving-guide,1611-4.html>, 2007.
- [2] Wikipedia, *Dynamic random-access memory*, http://en.wikipedia.org/wiki/Dynamic_random-access_memory, 2012.
- [3] Wikipedia, *Comparison of SSD with HDD*, http://en.wikipedia.org/wiki/Solid-state_drive#Comparison_of_SSD_with_hard_disk_drives, 2012.